

# Data Clorox CS513 Data Cleansing

Final Project Report

August 4, 2020

Asad Bin Imtiaz (Net ID: **aimtiaz2**), Muhammad Rafay (Net ID: **mrafay2**)

## 1 Introduction

This report summarizes the methods and tools, as well as the analysis and findings carried out for data wrangling, standardization and provenance workflow for the [1] *The New York Public Library* (NYPL), *What's on the menu?* Dataset. The dataset can be downloaded from the NYPL GitHub website [2]. The analysis and findings are part of Final Project for CS513: Theory and Practice of Data Cleaning Course from the university of Illinois, wherein the NYPL dataset [2] was used to create an end-to-end data wrangling and provenance workflow, together with data landscape analysis and findings, using data provenance and data cleansing techniques learned in the class. The goal of this project was to use several open-source tools and libraries for Data Wrangling and Data Provenance to come up with Data Cleaning Workflow which effectively cleans the selected dataset to high-quality standards with all the lineage and audit tracking available.

### 1.1 Tools and libraries

Following tools were used in this report:

- Python 3 with Jupyter notebook
- OpenRefine data cleaning tool [3]
- SQLite [4] with DB-Visualizer Pro 9.2 [5]
- Yes Workflow [4]
- Teradata 14 DWH [7]

### 1.2 Dataset

The New York Public library (NYPL) maintains a large collection of Menus (~45K) in their 'What's on the Menu' [2] dataset, which is openly available to download [1]. The dataset consists of CSV files with entities such as dish-by-dish menus from a variety of businesses from as early as 1850, and are used by historians, nutritionists and researchers around the globe to understand the patterns and to answer specific questions. The data is collected by taking photographs of menus over several years by volunteers and was digitized in the dataset form in NYPL Digital Gallery [1].

As with all the crowd-sourced gathered data, there are several gaps and inconsistencies in the data, as well as areas with potential for improvement in terms of the data formats, linking & lineage and its schema. The goal of this project is to identify the issues and fix them, keeping the provenance and transformation lineage to understand the cleansing workflow and to later reproduce the cleaned dataset on newer dataset versions.

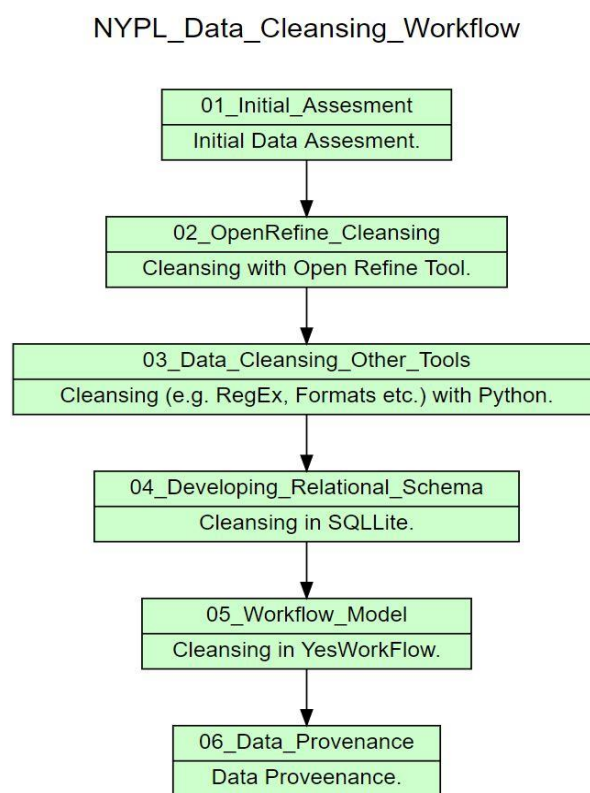
The initial assessment of data quality and respective issues are presented in chapter **Error! Reference source not found.** in detail.

### 1.3 Approach

The project work was divided into multiple tasks. Below is the task breakdown:

- Overview and initial assessment of the dataset.
- Data cleaning with Open-Refine [3]
- Data cleaning with other tools
- Developing a relational schema
- Creating a workflow model
- Developing provenance

Following is the high-level workflow illustration for the above tasks:



Each of the subtasks is discussed as a separate chapter in the following.

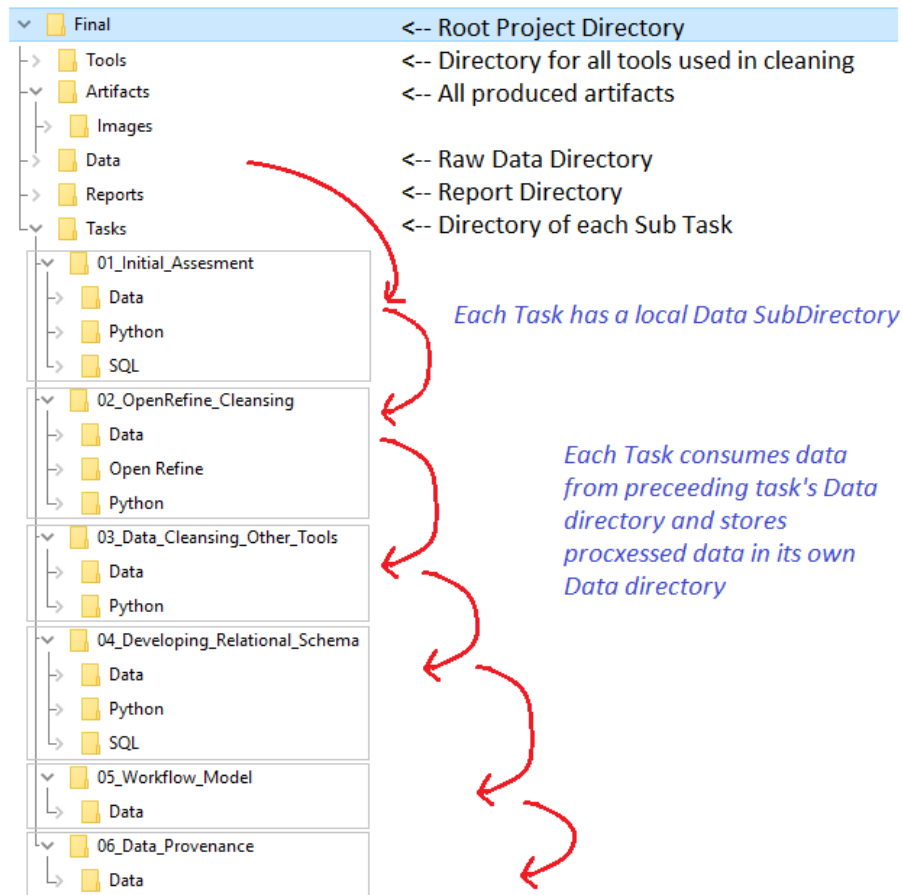
### 1.4 Project Setup

The project directories are set up hierarchically. The raw data is present in the root project directory in a subdirectory called **Data/**. The **Artifacts/** directory in the root project directory is for any artifacts, such as images generated by the data wrangling tasks. The **Tools/** directory contains executable tools used for data cleaning (such as SQLite) so that workflow could be started without any installations.

Each of the above 6 tasks has a subdirectory in the **Tasks/** directory. The project is set up in a directory structure where each of the above 6 tasks has their **Data/** directories, under the **Tasks/<Task\_Name>/** directory in the project root directory. The First task, takes the data from the Raw Data directory in the root project folder (**<Root>/Data/**) and stores processed

data in its local directory (**<Root>/Tasks/<Task\_Name>/Data/**). Each subsequent task takes the data from the preceding task's Data directory and stores it in its own Data directory after processing. This directory structure is structured so that retrospective provenance of data flow can be maintained for each step.

Below is how the project directory setup is built.



The **workflow.sh** or **workflow-annotated.sh** are the main workflow scripts that initiate the entire cleaning pipeline.

The project repository is located on Git-Hub [8] at:

[https://github.com/AsadBinImtiaz/CS513\\_Data\\_Cleaning\\_Final](https://github.com/AsadBinImtiaz/CS513_Data_Cleaning_Final)

SHA-1 Key (Read-Only) to access the repository can be found in appendix of this report.

## 2 Overview and initial assessment of the dataset

In the following subsection, the structure and content on the dataset are inspected before starting with the data wrangling and provenance workflow, to get familiarity with data schema and a feel for apparent data quality issues present in the data. There may be more issues in data that would be discussed in subsequent chapters with corresponding tasks.

The initial assessment was performed to get an understanding of the data quality in general and to identify methodology and tools for subsequent tasks. Scoped in this task was also an exploration of the data structures and types, and an understanding of the relationship among these entities.

In the following chapters is the description of the data and data objects, as an outcome of the initial assessment.

## 2.1 Data Structure

The entire dataset consists of four character-delimited files described below:

### ***I. Dish.csv***

This file contains all dishes with their dish names listed on the menu along with their respective pricing and chronology information. Each record represents a specific dish offered by a business and listed on the menu. Each dish has an identifier that uniquely identifies it and is referenced as a foreign key on other entities. The file may be considered as dimension entity for dishes.

### ***II. MenuItem.csv***

This file contains menu items that link a menu page entity with dish entities as foreign references. Each record is identified by a unique identifier and carries other information such as associated dish price and x/y position of the image of the menu page. Each menu item refers to a dish in the dish entity and a menu page in a menupage entity, thereby creating a link between a dish and a menu page.

### ***III. MenuPage.csv***

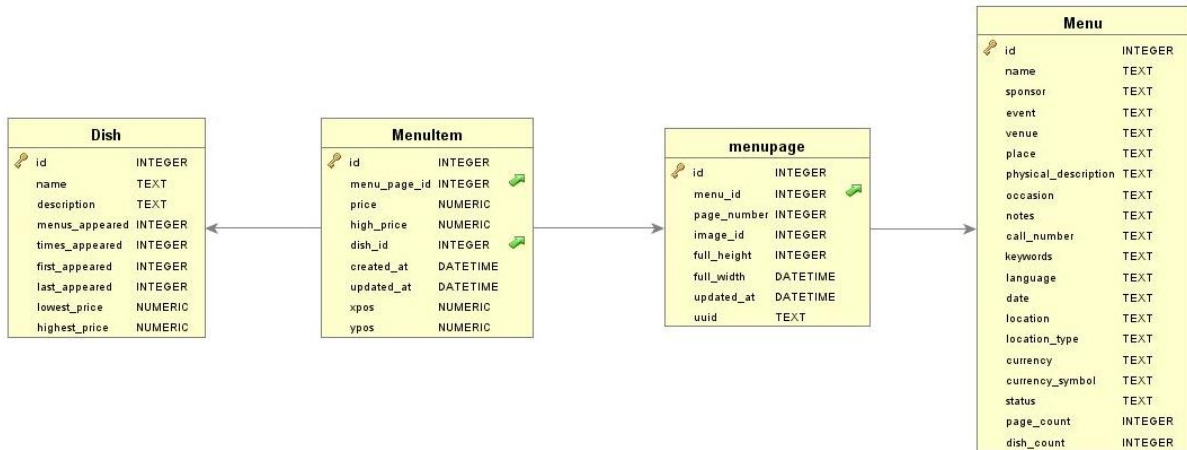
This file contains menu page records. Each item is identified by a unique identifier and links a menu item with a menu. Additional information such as page photo image number and page dimensions also appear here. Every record keeps references for the menu item identifier and menu identifier to link these entities together.

### ***IV. Menu.csv***

This file contains all individual menus, each associated with a unique id. Each menu has an identifier that uniquely identifies it and is referenced as a foreign key on other entities. Associated data includes the occasion, venue, and event information and chronological information such as created and updated dates and times. Other important fields present in this file include the location where the menu is offered, the associated currency in use for menu items, the language, and the status of the menu.

## 2.2 Data Schema

The raw data were imported in an SQL-Lite instance and visualized using the DB-Visualizer tool. The ER diagram generated from DB-Visualizer is shown in the figure below:



The cardinalities of objects with respect to one another are found to be (only list for entities among where a direct link is possible):

From Entity	To Entity	Cardinality of relation
Dish	MenuItem	1:N
MenuItem	Dish	1:1
MenuItem	MenuPage	1:1
MenuPage	MenuItem	1:N
MenuPage	Menu	1:1
Menu	MenuPage	1:N

Given the cardinalities above and the Initial Quality Assessment of the dataset in section 2.4, we have assessed that some rows are duplicate in the data and they need to be merged according to certain criteria.

**e.g.** in Dish.csv, the following dish has a duplicate due to difference in cases

```
Dish.id = 1, Dish.name = 'Consomme printaniere royal'
Dish.id = 397198, Dish.name = 'Consomme Printaniere Royal'
```

Although the cardinality from MenuItem to Dish of 1:1 is maintained due to the unique Dish.id but just due to differences in cases for dish names, there should not exist multiple records representing the same dish. These problems can only be analyzed and mitigated after the initial cleansing.

## 2.3 Data Types

The diagram shows entities and links for data objects present in the data set. Most of the raw data was imported as strings of characters. However, the initial assessment showed the following data types for the fields:

Entity: Dish					
Field Name	Type	Precision	Format	Key	Null
Id	Integer		(10)9	PK	N
Name	String	1387	X(1387) Unicode		N
Description	String	0	X(1)		Y

Menus_appeared	Integer		-(10)9		N
Times_Appeared	Integer		-(10)9		N
First_Appeared	Integer		(4)9		N
Last_Appeared	Integer		(4)9		N
Lowest_Price	Numeric	2	-----.99		Y
Highest_Price	Numeric	2	-----.99		Y

Entity: MenuItem					
Field Name	Type	Precision	Format	Key	Null
Id	Integer		(10)9	PK	N
Menu_Page_Id	Integer		(10)9	FK	N
Price	Numeric	2	----.99		Y
High_price	Numeric	2	----.99		Y
Dish_id	Integer		(10)9	FK	Y
Created_at	Timestamp p(0) With zone		YYYY-MM-DD hh:mm:ss(0) Z		N
Updated_at	Timestamp p(0) With zone		YYYY-MM-DD hh:mm:ss(0) Z		N
Xpos	Numeric	6	-.999999		N
Ypos	Numeric	6	-.999999		N

Entity: MenuPage					
Field Name	Type	Precision	Format	Key	Null
Id	Integer		(10)9	PK	N
Menu_Id	Integer		(10)9	FK	N
Page_Number	Integer		----.99		Y
Image_Id	String	15	X(15)		N
Full_height	Integer		(4)9		Y
Full_width	Integer		(4)9		Y
Updated_at	String	36	X(36) [UUID]		Y
Uuid	Numeric	2	-.999999		Y

Entity: Menu					
Field Name	Type	Precision	Format	Key	Null
Id	Integer		(10)9	PK	N
name	String		(10)9	FK	N
sponsor	String		----.99		Y
event	String	15	X(15)		Y
venue	String		(4)9		Y
place	String		(4)9		Y
physical_description	String	36	X(36) [UUID]		Y
occasion	String	2	-.999999		Y
notes	String		X(260)		Y
call_number	String		X(40)		Y
keywords	String		X(0)		Y
language	String		X(0)		Y
date	Date		YYYY-MM-HH		Y
location	String		X(0)		Y
location_type	String		X(127)		Y
currency	String		X(26)		Y
currency_symbol	String		X(4)		Y
status	String		X(9)		Y
page_count	Integer		(10)9		Y
dish_count	Integer		(10)9		Y

## 2.4 Data Quality initial assessment

After an initial assessment of data, various data quality issues were apparent. The following is a small summary of issues discovered at the initial assessment. following data quality checks were performed and respective violations were listed.

### **Dish.csv file:**

In the dish file, the names were not stored in the standard case. Some names are upper-case, some lower-case and other mixed-case. The names also contained extra spaces and invalid characters. The description field was empty. Several dishes had menus\_appeared or times\_appeared value as 0. The correctness of these fields was also not correct; several dishes times\_appeared value less than menus appeared. Similar issues were found with first\_appeared and last\_appeared, with many dishes having the first appearance earlier than the menu date or last appearance year earlier than the first appearance year. there were missing values in prices which may be filled or corrected with prices in menu\_item.

### **Menuitem.csv file:**

In the Menuitem file, several menus had missing references to dishes. These were referential integrity issues such as a dish reference not existing as a dish in the dish file. Many prices were missing and could be completed from the dish file. The same issues were present for the highest price field. In many cases, the Highest price was less than the lowest price. Similarly, in many cases, updated\_at time-stamp is earlier than created\_at time-stamp.

### **MenuPage.csv file:**

In the MenuPage file, there were referential integrity issues. Several menu ids referenced were non-existent in the menu file. The page numbers were also not correct and sometimes had negative values. The Image id field was also not standard and had alphanumeric values in a few cases.

### **Menu.csv file:**

In the Menu file, there were several issues with textual fields, such as missing values, non-standard cases, having invalid characters, etc. There were values such as '?' or '[restaurant and/or location not given.]' with the same semantic as missing values. Several fields such as keywords, languages, and location\_type were empty with no value in the entire data file. Other fields such as status contained only a single value all the time. The dates were also not in a standard and consistent format. The dish count value in the menu in a few cases was not correct if associated dishes were counted. The call numbers, mostly numeric, in few cases, had trailing alpha characters

## 2.5 Use cases

The usefulness of data can be judged by the potential use case it may serve. The raw data has many issues as listed above but may still be valuable for several analytic scenarios and use cases. However once cleaned, further uses can be foreseen, some of them listed below:

### **2.5.1 Fitness for use as is**

Although the data in raw form is not clean and has several anomalies, it may still serve several use cases, including but not limited to:

- The data may be used by business owners to generate menus and see historical variations in the dish listings.
- The data may be used to have an estimate on dish popularity based on the number of times a dish is listed on all menus. It may also be used to find dishes previously listed but not offered anymore as a criterion of the unpopularity of the dish.
- One can analyze how menus have changed across time in terms of their page size and number of dishes.

### **2.5.2 Fitness for use after data wrangling**

Once the initial quality issues have been addressed the data will be fit for the following and other similar use cases:

- Once the correct chronology of menu listings may be established, the dataset may be used to study how eating preferences evolved by correlating them with the popularity of dishes for a given interval.
- Once duplication in dishes is merged, upselling, cross-selling, and competitive pricing analysis for dishes across locations may be performed.
- If a review dataset (such as yelp reviews) may be combined with this data (based on cleansed business names etc.), rating or sentiment analysis of dishes may be performed per restaurant.

### **2.5.3 Unfit for use even after data wrangling**

- The data is unfit for rating or performing sentiment analysis of the dishes on its own even after data cleaning without combining it with a review's dataset.
- It is not possible to verify the validity of dish characteristics or their listing, as well as the completeness of the Menus with respect to their images, due to unavailability of this auxiliary data in the raw dataset.

## **3 Data cleaning with Open-Refine (and other tools)**

The data was cleaned with Open-Refine Tool (formerly known as Google-Refine) and other Programming tools such as python and R. In the following these cleaning tasks are mentioned in detail.

### **3.1 Identifying Data Cleaning Steps for the use cases**

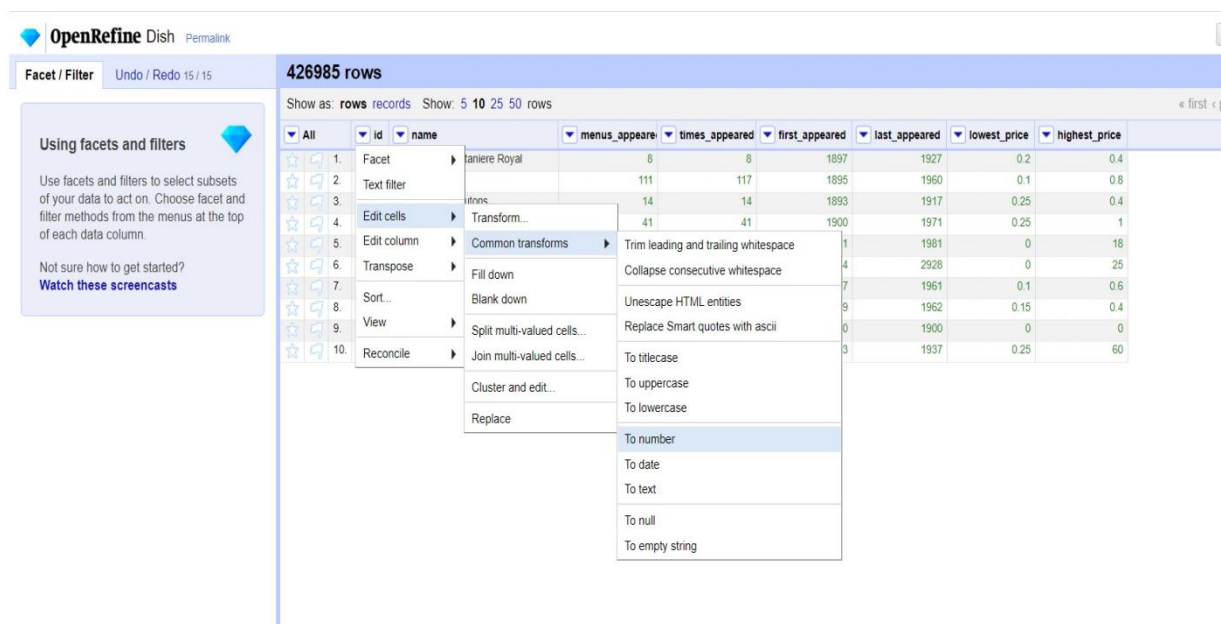
To make the data fit for the above-mentioned use cases, data was needed to be cleaned and standardized. Several cases of missing values in records, non-standard formats, and cases, invalid characters, additional white spaces, inconsistency in spellings of textual, duplicates, and other issues involving completeness and correctness were identified in Initial Assessment. Many of these were fixed with the OpenRefine tool. Moreover, there existed violations such as uniqueness and referential integrity among the data entries, but these addressed in Section 4 with a relational database.



OpenRefine allowed us to facet, filter, transform, and cluster data. It also allows operations on data using the GREL language. All features available under a GUI makes it easy to learn and use. It also keeps a track of data provenance. But still, it has its limitations so further cleaning process using completed python language for the tasks that we won't be able to do using OpenRefine such as regular expression matching. We will be using the pandas library which is a very fast data analysis and manipulation tool in python.

### 3.2 Data cleaning with OpenRefine

Open Refine is an open-source tool for data wrangling. It reads the CSV files as data tables and performs several data cleaning tasks such as type casting, case standardization, and clustering. OpenRefine was used to cleanse cases of format standardization, for conversions of type and clustering base of same or similar values together in textual data.



**Note:** All file were read with encoding= 'utf-8'

### 3.3 Open Refine Data Cleaning Process

#### Dish.csv file:

There were several issues identified which were suitable for cleansing with OpenRefine. The below table contains a summary of these issues and their rectification in OpenRefine.

Field(s)	Issue type	Issue Description	Resolution
ID menus_appeared times_appeared first_appeared last_appeared lowest_price highest_price	Type Cast	Convert to Number	Text transformed with <b>toNumber()</b> <b>426985</b> Records affected

Description	Completeness	The entire field is empty and unusable	1 Column removed
Name	Extra Spaces	There is extra space <ul style="list-style-type: none"> <li>Leading dish names</li> <li>Training dish names</li> <li>In between dish names</li> </ul> <u>e.g.</u> Dish.id = 131274, Dish.name = 'Consomme printaniere royal'	Text transformed with <b>trim()</b> <b>9226 cells updated</b>
Name	Extra Spaces	There is extra space <ul style="list-style-type: none"> <li>Leading dish names</li> <li>Training dish names</li> <li>In between dish names</li> </ul> <u>e.g.</u> Dish.id = 131274, Dish.name = 'Consomme printaniere royal' Dish.id = 397198, Dish.name = ' " " kidneys'	Text transformed with <b>replace(/s+/, ' ')</b> <b>6554 cells updated</b>
Name	Case Standardization	The name does not appear in the standard case. Some names are upper-case, some lower-case and other mixed-case. <ul style="list-style-type: none"> <li>There are 426985 distinct dishes</li> <li>There are 398443 distinct dishes case insensitive</li> </ul>	Text transformed to title case with <b>toTitlecase()</b> <b>284173 cells updated</b>
Name	Clustering	Similar names to be clustered.	Clustered 4006 texts with <b>MassEdit</b>
Name	Invalid characters	There are invalid characters like (!, @, #, { etc.) in dish names <u>e.g.</u> Dish.id = 2839, Dish.name = 'E. & J. B. ***'	Removed invalid characters with <b>GREL</b> <b>6554 cells updated</b>

In the project directory, **Tasks/02\_OpenRefine\_Cleansing/Open Refine/** the audit log for Dish table can be found in **Dish.json**

### **Menuitem.csv file:**

Menu Items did not have a lot of text fields, therefore much of the cleansing was done later with SQL in a database. The below table contains a summary of issues identified and their rectification in OpenRefine.

Field(s)	Issue type	Issue Description	Resolution
Id menu_page_id price high_price dish_id xpos ypos	Type Cast	Convert to Number	Text transformed with <b>toNumber()</b> <b>1334419 Records updated</b>
created_at updated_at	Type Cast	Convert to Date in format YYYY-MM-DD HH:MI:SS without time zone	<b>grel:value.split('UTC')[0]and toDate()</b> <b>1334419 Records updated</b>

In the project directory, **Tasks/02\_OpenRefine\_Cleansing/Open Refine/** the audit log for Menu Items table can be found in **MenuItem.json**

### **MenuPage.csv file:**

Menu page records also did not have a lot of text fields, therefore much of the cleansing was done later with SQL in a database. The below table contains a summary of issues identified and their rectification in OpenRefine.

Field(s)	Issue type	Issue Description	Resolution
Id menu_id Page_number full_height full_width	Type Cast	Convert to Number	Text transformed with <b>toNumber()</b> <b>66937</b> Records updated
image_id	Correctness transform	remove leading/trailing characters from image number	Text transform with <b>grel:value.match(/(\w+)(\d+)/)</b> <b>23</b> cells updated
image_id	Type Cast	Convert to Number	Text transformed with <b>toNumber()</b> <b>66608</b> Records updated

In the project directory, **Tasks/02\_OpenRefine\_Cleansing/Open Refine/** the audit log for Menu page table can be found in **MenuPage.json**

### **Menu.csv file:**

The menu file contained a lot of text fields, ideal for cleansing with the OpenRefine tool. It had a lot of potential for clustering and pattern matching for the cleansing of fields. The below table contains a summary of issues identified and their rectification in OpenRefine.

Field(s)	Issue type	Issue Description	Resolution
Id Page_Count Dish_Count	Type Cast	Convert to Number	Text transformed with <b>toNumber()</b>
Keywords Language location_Type	Completeness Empty fields	The entire fields are empty and unusable	<b>3 Columns removed</b>
Name	Standardization	The values do not appear in the standard case. Some values are upper-case, some lower-case, and other mixed-case.	Text transformed to title case with <b>toUppercase()</b>
Name	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> and <b>replace(/\s+/, ' ')</b> <b>9+1 cells updated</b>
Name	Clustering	Similar names to be clustered.	Clustered 588 texts with <b>MassEdit</b>
Sponsor	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> and <b>replace(/\s+/, ' ')</b> <b>3+6 cells updated</b>

Sponsor	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toUppercase()</b>
Sponsor	Clustering	Similar names to be clustered.	Clustered 899 cells with <b>MassEdit</b>
Event	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> and <b>replace(Λs+/, ' ')</b> <b>3+6 cells updated</b>
Event	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toUppercase()</b>
Event	Clustering	Similar names to be clustered.	Clustered 5314 cells with <b>MassEdit</b>
Place	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> <b>45 cells updated</b>
Place	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toUppercase()</b> <b>899 cells updated</b>
Place	Clustering	Similar names to be clustered.	Clustered 3184 cells with <b>MassEdit</b>
Physical_Description	Correctness	Remove Extra spaces	Text transformed to title case with <b>replace(Λs+/, ' ')</b> <b>38 cells updated</b>
Occasion	Correctness	Remove Extra spaces	Text transformed to title case with <b>replace(Λs+/, ' ')</b> <b>3 cells updated</b>
Occasion	Clustering	Similar names to be clustered.	Clustered 2779 cells with <b>MassEdit</b>
Occasion	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toUppercase()</b> <b>17 cells updated</b>
Notes	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> and <b>replace(Λs+/, ' ')</b> <b>125+195 cells updated</b>
Notes	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toUppercase()</b> <b>3873 cells updated</b>
Notes	Clustering	Similar names to be clustered.	Clustered 1355 cells with <b>MassEdit</b>
Call_Number	Truncation	remove extra suffix characters	Text transformed with <b>grel: value.split(' ')[0]</b> and <b>value.match(/(Λd+-Λd+)/)[0]</b> <b>1093+22 cells updated</b>
Location	Correctness	Remove Extra spaces	Text transformed to title case with <b>trim()</b> and <b>replace(Λs+/, ' ')</b> <b>14+555 cells updated</b>
Location	Clustering	Similar names to be clustered.	Clustered cells with <b>MassEdit</b> <b>4205+5+48+24+36 cells updated</b>
Currency	Standardization	The values do not appear in the standard case.	Text transformed to title case with <b>toTitlecase()</b> <b>118 cells updated</b>
Currency	Correctness	Remove invalid characters	Text transformed with <b>grel: value.split('(')[0]</b> <b>43 cells updated</b>
Location venue	Clustering	Clustering with the mass edit	<b>1283+1463+5104 cells updated</b> <b>2243+21 cells updated</b>

name			609+55 cells updated
event			2325 cells updated
occasion			270+2 cells updated

In the project directory, **Tasks/02\_OpenRefine\_Cleansing/Open Refine/** the audit log for Menu page table can be found in **Menu.json**

### 3.4 Data cleaning with Pandas (Python)

Since the title case in OpenRefine does not ignore brackets, e.g for text “hello (world)”, OpenRefine would TitleCase it to “Hello (world)” and not to “Hello (World)”, most case conversions in OpenRefine were limited to UpperCase. Python is intelligent in such conversions, and reconverted appropriately, in our example’s case to “Hello (World)”.

With python, the following corrections were performed:

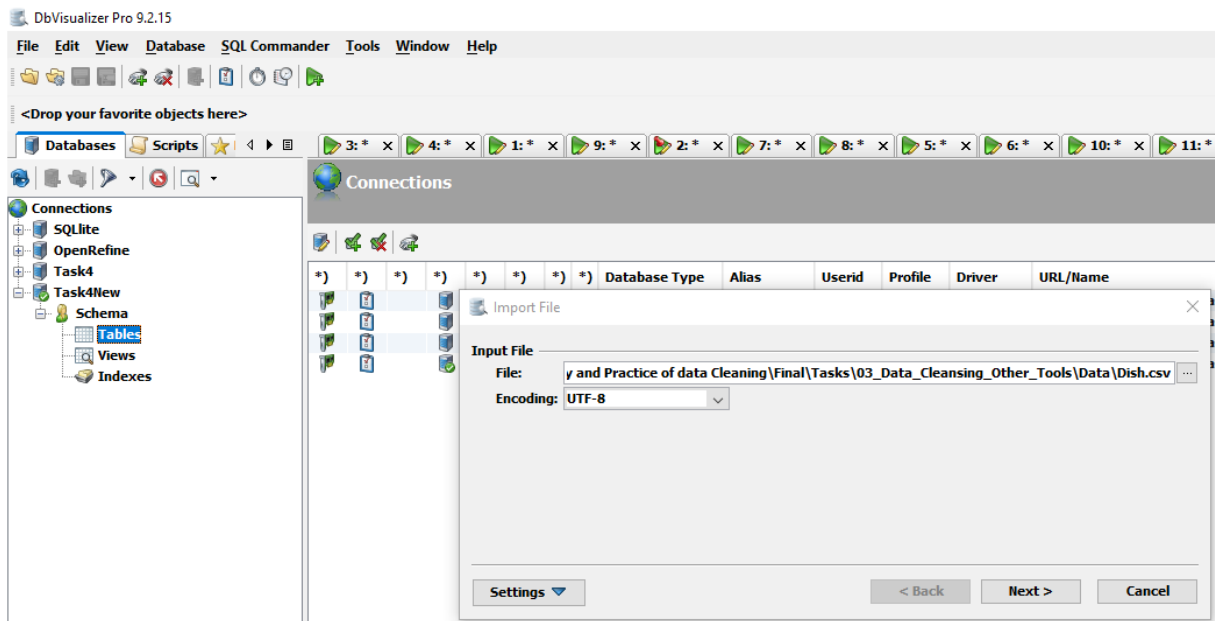
File	Field(s)	Issue Description	Resolution
Dish.csv	name	Convert To TitleCase	<code>dish_df.name.str.title()</code>
Dish.csv	name	Replace & with And	<code>dish_df.name.str.replace(' &amp; ',' And ')</code>
Dish.csv	lowest_price highest_price	Set precision to 2 decimal places	<code>float_format='%.2f'</code>
Menu.csv	name sponsor event venue place physical_description occasion notes location	Convert To TitleCase	<code>str.title()</code>
Menu.csv	Date	Correct Date Format	<code>Format = 'YYYY-MM-DD'</code>
Menu.csv	name sponsor event venue place physical_description occasion notes location	Replace & with And	<code>dish_df.name.str.replace(' &amp; ',' And ')</code>
MenuItem.csv	xpos ypos	Set precision to 6 decimal places	<code>float_format='%.6f'</code>
MenuItem.csv	created_at updated_at	Correct Timestamp Format	<code>Format = 'YYYY-MM-DD MI:HH:SS'</code>
MenuItem.csv	price high_price	Set precision to 2 decimal places	<code>float_format='%.2f'</code>
MenuPage.csv	image_id	remove suffixes:	<code>str.replace('psnypl_rbk_', '').replace('ps_rbk_', '')</code>

## 4 Developing a relational schema

In this sub-task, the data cleansed with OpenRefine and Python was loaded into an SQL-Lite DB and further cleansing was performed with Structured Query Language (SQL). The following sub-chapters discuss the steps and processes taken to cleanse the data.

## 4.1 Loading data in SQL-Lite Database

The four CSV files were loaded into an SQL-lite instance via DB visualizer (v9.2.15 pro) using its internal SQL-Lite connector (v3.8).



Once the data was inside the SQL-Lite DB, SQL queries were written to identify issues relating to missing values, domain type conversion, and integrity constraints. These subtasks are discussed in the following chapters.

## 4.2 Identifying Remaining issues & referential integrity violations

Within the imported data in SQL-Tables, the following issues were identified.

### Table Dish

Field	Issue type	Description
Name	Invalid characters	There are invalid characters like (!, @, #, { etc.) in dish names e.g. Dish.id = 2839, Dish.name = 'E. & J. B. ***'
Menus_appeared	Plausibility	There are 2412 Dishes with menus_appeared = 0
Menus_appeared	Correctness	There are differences in menus appeared and actual menu count for dish e.g. id = 19, menu_appeared = 16, actually appeared = 15
Times_appeared	Plausibility	Several 0 or negative values -- 1 Dishes appeared -10 times ??? [MIN] -- 11900 dishes appeared 0 times !!! -- 372 dishes appeared 19 Menus [MAX]
Times_appeared	Correctness	There are differences in times appeared and actual count for the dish in menus e.g.

		id =17, times_appeared = 535, actually appeared = 536
First_appeared	Plausibility	Many dishes have first appearance earlier than menu date or later than last_appeared year
Lowest_price	Correctness	Dish lowest price should not be negative. Several dishes have 0 lowest prices. It may not be an issue but worth analyzing especially when nulls are allowed.
Highest_price	Correctness	Dish's highest price should not be negative.
Name	Duplication	Same standardized dish name has multiple entries with different context

#### **Table MenuItem:**

Field	Issue type	Description
dish_id	Null as FK, lineage has broken	241 menu items have no value for dish id e.g. menu_item.id = 19171 , Menu_item.dish_id = NULL
dish_id	Referential integrity	3 dish ids in menu item which do not exist in the dish.csv e.g. menu_item.id = 619133 , Menu_item.dish_id = 220797
Price	Completeness	More than 446K menus have null in price. It may be overwritten by an average dish price from <ul style="list-style-type: none"> <li>The highest price in menu item (58 cases)</li> <li>Menu Items for the same dish</li> <li>Dish lowest and highest prices</li> </ul>
high_price	Completeness	More than 1.2M menu items have no high price. It may be overwritten from <ul style="list-style-type: none"> <li>Lowest price in menu item (~800k cases)</li> <li>Corresponding dish highest prices</li> </ul>
Price	Correctness	in 1278 cases, High_price is strictly less than the price e.g. MenuItem.id = 1455, price = 40, high_price = 0.4
Created_at	plausibility	in 2874 cases updated_at timestamp is earlier than created_at timestamp

#### **Table MenuPage**

Field	Issue type	Description
id	Referential integrity	40334 Menu pages are not referred to by any menu items. This may not be a problem (e.g. title page etc.) but worth analyzing.
menu_id	Referential integrity	5803 Menu ids in menu page which do not exist in Menu file e.g. menu_page.id = 119 , Menu_Page.Menu_id = 12460
page_number	Plausibility	Range of Page_number is (1,74) Is 74 pages long menu plausible?
page_number	Completeness	Missing 1202 values. Can be constructed from image_id and menu_id
Image_id	Correctness	in 92 cases, menupage refers to the same menu_id and same page_number but different image_id
full_height	Completeness	in 329 cases, full_height is null but image_id is known

full_width	Completeness	in 329 cases, full_width is null but image_id is known
uuid	Uniqueness	Having UUID in it, this field has 2922 duplicated ids, several ids repeating as many as 10 times.

### Table Menu

Field	Issue type	Description
name	Consistency	There are values such as <i>[not given]</i> or <i>[restaurant name and/or location not given]</i> . Do these values sound reasonable when nulls/blanks are allowed and vice versa?
sponsor	Consistency	57 records have '?' as value. 30 records have '[restaurant and/or location not given.]' Are these reasonable when nulls/blanks are allowed and vice versa?
event	Consistency	Similar or same values e.g dinner, dinner/dinner, [dinner], daily dinner, (dinner) ?, <Blank>
date	Correctness	Few Invalid Dates e.g 2928-03-26 586/17545 Missing Values
dish_count	Correctness	in 214 cases, the value of dish count is different to distinct dishes the menu can be connected to

## 4.3 SQL queries to check and meet the integrity constraints

### Table Dish

The first correction which was performed was a truncation of dish names longer than 500 characters. 104 such dishes were identified and realized to have descriptions along with names. These were manually corrected.


All dish records where last\_appeared as earlier than first\_appeared were corrected.

```


110 UPDATE Dish
111 SET last_appeared = first_appeared
112 WHERE last_appeared < first_appeared
113 ;
114

```

110:2 [12983] INS

 Log

☒ Preprocess script
☒ Log to GUI
☐ Log to File

 15:33:58 [UPDATE - 0 rows, 0.679 secs] Command processed. No rows were affected  
... 1 statement(s) executed, 0 rows affected, exec/fetch time: 0.679/0.000 sec [1 successful, 0 errors]

Moreover, any first or last appearance of a dish later than the current year was adjusted to the current year:




```



119 UPDATE Dish
120 SET last_appeared = CAST(EXTRACT( YEAR FROM CURRENT_DATE) AS INTEGER)
121 WHERE last_appeared > CAST(EXTRACT( YEAR FROM CURRENT_DATE) AS INTEGER)
122 ;
123
124 UPDATE Dish
125 SET first_appeared = CAST(EXTRACT( YEAR FROM CURRENT_DATE) AS INTEGER)
126 WHERE first_appeared > CAST(EXTRACT( YEAR FROM CURRENT_DATE) AS INTEGER)
127 ;

```

119:12 [13325] INS

 Log

☒ Preprocess script ☒ Log to GUI ☐ Log to File

 15:40:21 [UPDATE - 179 rows, 0.873 secs] Command processed  
 15:40:22 [UPDATE - 11 rows, 0.730 secs] Command processed  
 ... 2 statement(s) executed, 190 rows affected, exec/fetch time: 1.603/0.000 sec [2 successful, 0 errors]


All values for dishes that have the same name were consolidated into single records and extra records were deleted.

```

357 DELETE FROM Dish WHERE Id IN
358 (SELECT old_id FROM Dish_Temp WHERE old_id <> new_id)

```

345:2 [18959] INS

 Log

☒ Preprocess script ☒ Log to GUI ☐ Log to File

 16:28:41 [CREATE - 0 rows, 1.092 secs] Command processed. No rows were affected  
 16:28:42 [UPDATE - 746000 rows, 0.790 secs] Command processed  
 16:28:42 [DELETE - 36701 rows, 0.172 secs] Command processed  
 16:28:42 [UPDATE - 23371 rows, 0.403 secs] Command processed

### Table Menu


The text which represents missing values such as ?,[Not Given], unknown, etc were converted to unique singletons.

```

172 UPDATE Menu
173 SET Sponsor = '[Not Given]'
174 WHERE 1=0
175 OR TRIM(sponsor) = '?'
176 OR TRIM(sponsor) = ''
177 ;
178
179 UPDATE Menu
180 SET NAME = '[Not Given]'
181 WHERE 1=0
182 OR TRIM(NAME) = '?'
183 OR TRIM(NAME) = ''
184 ;
185
186 UPDATE Menu
187 SET Sponsor = NAME
188 WHERE NAME IS NOT NULL
189 AND TRIM(sponsor) = '[Not Given]'
190 AND TRIM(NAME) <> '[Not Given]'
191 ;
192
193 UPDATE Menu
194 SET NAME = Sponsor
195 WHERE sponsor IS NOT NULL
196 AND COALESCE(TRIM(NAME), '[Not Given]') = '[Not Given]'
197 AND TRIM(sponsor) <> '[Not Given]'
198 ;
199

```

172:12 [15149] INS



☒ Preprocess script ☒ Log to GUI ☐ Log to File

```

16:12:18 [UPDATE - 57 rows, 0.039 secs] Command processed
16:12:18 [UPDATE - 0 rows, 0.045 secs] Command processed. No rows were affected
16:12:18 [UPDATE - 134 rows, 0.084 secs] Command processed
16:12:19 [UPDATE - 5270 rows, 0.050 secs] Command processed
... 4 statement(s) executed, 5461 rows affected, exec/fetch time: 0.218/0.000 sec [4 successful, 0 errors]

```

### Table MenuPage

All page numbers in the Menu page which were negative or had gaps or overlaps were resequenced again with image\_ids.

```

261 UPDATE MenuPage
262 FROM
263 (
264     SELECT id Src_Id, page_number, ROW_NUMBER() OVER (PARTITION BY menu_id ORDER BY id) new_page_id
265     FROM MenuPage x WHERE COALESCE(page_number, -1) = -1
266 ) X
267 SET Page_number = X.new_page_id
268 WHERE id = X.Src_Id
269 ;

```

All menu-pages not referencing menus, and all menu-items not referencing dishes or menus were iteratively removed.

### Table MenuItem:

Records, where high\_price was less than price or created timestamp, was later than updated timestamp were corrected.

```

296 UPDATE MenuItem
297 SET high_price = price
298 WHERE price > high_price
299 ;
300
301 UPDATE MenuItem
302 SET updated_at = created_at
303 WHERE created_at > updated_at
304 ;
305
306 UPDATE MenuItem
307 SET high_price = price
308 WHERE high_price IS NULL
309 AND price IS NOT NULL
310 ;

```

294:66 [17473] INS

Log

☒ Preprocess script ☒ Log to GUI ☐ Log to File

```

16:25:33 [UPDATE - 1278 rows, 0.383 secs] Command processed
16:25:34 [UPDATE - 2874 rows, 0.087 secs] Command processed
16:25:34 [UPDATE - 796337 rows, 0.636 secs] Command processed
... 3 statement(s) executed, 800489 rows affected, exec/fetch time: 1.106/0.000 sec [3 successful, 0 errors]

```

### Fixing RI:

All records were removed where referenced id did not exist in the referenced table.

```

385 DELETE FROM MenuPage WHERE menu_id IN (SELECT id FROM Menu);
386 DELETE FROM MenuItem WHERE menu_page_id NOT IN (SELECT id FROM MenuPage);
387 DELETE FROM MenuItem WHERE dish_id NOT IN (SELECT id FROM dish);

```

The values of times\_appeared in the dish table were correct by counting the number of its appearances in menus. Moreover, the value for menus\_appeared was corrected by associating distinct menus.

```

5 UPDATE MenuItem
6 FROM ( SELECT M.id Src_Id, M.high_price, M.Dish_Id, Max_Price FROM MenuItem M
7       JOIN ( SELECT dish_id, MAX(price) Max_Price FROM MenuItem WHERE price IS NOT NULL GROUP BY 1 ) D
8       ON M.Dish_Id = D.Dish_ID
9       WHERE 1=1 AND COALESCE(Max_Price,0) > 0 AND COALESCE(m.high_price,-1) = -1
10      GROUP BY 1,2,3,4
11 ) X
12 SET
13     high_price = Max_Price
14 WHERE id = Src_Id
15 ;

```

```

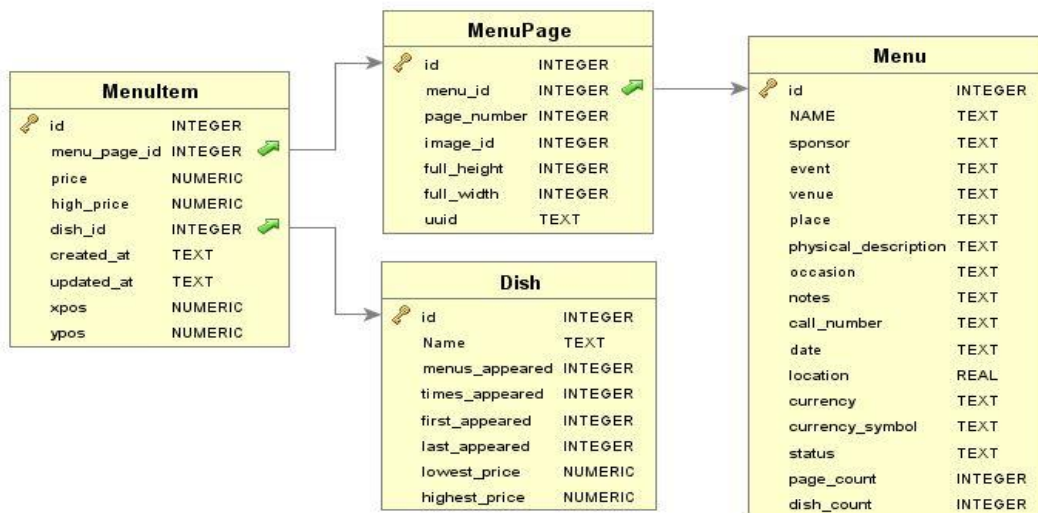
1
2 UPDATE MenuItem SET high_price = Price WHERE high_price < Price;
3 UPDATE MenuItem SET price = high_Price WHERE Price IS NULL AND high_price IS NOT NULL

```

Similarly first\_appeared and last\_appeared in the dish table was corrected from the year of the menu. The Dish prices were also adjusted from MenuItem and vice versa where missing.

## 4.4 Creating a New SQL Schema for cleaned data

After cleaning the data, a new schema was generated, as shown in the illustration below:



Below are the final Data Definition Language (DDL) definitions for RDBMS. The definitions show each field and its appropriate types and nullability.

#### Dish

```

CREATE SET TABLE Dish
(
    Id                INTEGER NOT NULL,
    Name              VARCHAR(500) CHARACTER SET UNICODE NOT CASESPECIFIC NOT NULL,
    menus_appeared    INTEGER NOT NULL,
    times_appeared    INTEGER NOT NULL,
    first_appeared    INTEGER NOT NULL,
    last_appeared     INTEGER NOT NULL,
    lowest_price       DECIMAL(18,2) NOT NULL,
    highest_price      DECIMAL(18,2) NOT NULL,
    PRIMARY KEY (id)
);
  
```

#### Menu

```

CREATE SET TABLE Menu
(
    Id                INTEGER NOT NULL,
    Name              VARCHAR(300) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Sponsor           VARCHAR(150) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Event             VARCHAR(200) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Venue             VARCHAR(50) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Place             VARCHAR(120) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Physical_description VARCHAR(120) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Occasion          VARCHAR(100) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Notes             VARCHAR(300) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Call_number       VARCHAR(10) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Date              DATE FORMAT 'YYYY-MM-DD',
    Location          VARCHAR(500) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Currency           VARCHAR(25) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Currency_symbol   VARCHAR(5) CHARACTER SET UNICODE NOT CASESPECIFIC,
    Status            VARCHAR(15) CHARACTER SET UNICODE NOT CASESPECIFIC NOT NULL,
    Page_count        INTEGER,
    Dish_count        INTEGER,
    PRIMARY KEY (id)
);
  
```

#### MenuItem

```

CREATE SET TABLE MenuItem
(
    Id                INTEGER NOT NULL,
    Menu_Page_Id      INTEGER NOT NULL,
    Price             DECIMAL(18,2),
    High_Price        DECIMAL(18,2),
    Dish_Id           INTEGER NOT NULL,
    Created_at        TIMESTAMP(0),
    Updated_at        TIMESTAMP(0),
    Xpos              DECIMAL(18,5),
    Ypos              DECIMAL(18,5),
    PRIMARY KEY (id),
    CONSTRAINT Dish_FK FOREIGN KEY (dish_id)
        REFERENCES Dish (id) ON UPDATE NO ACTION ON DELETE NO ACTION,
    CONSTRAINT MenuPage_fk1 FOREIGN KEY (menu_page_id)
        REFERENCES MenuPage (id) ON UPDATE NO ACTION ON DELETE NO ACTION
);
  
```

#### MenuPage

```

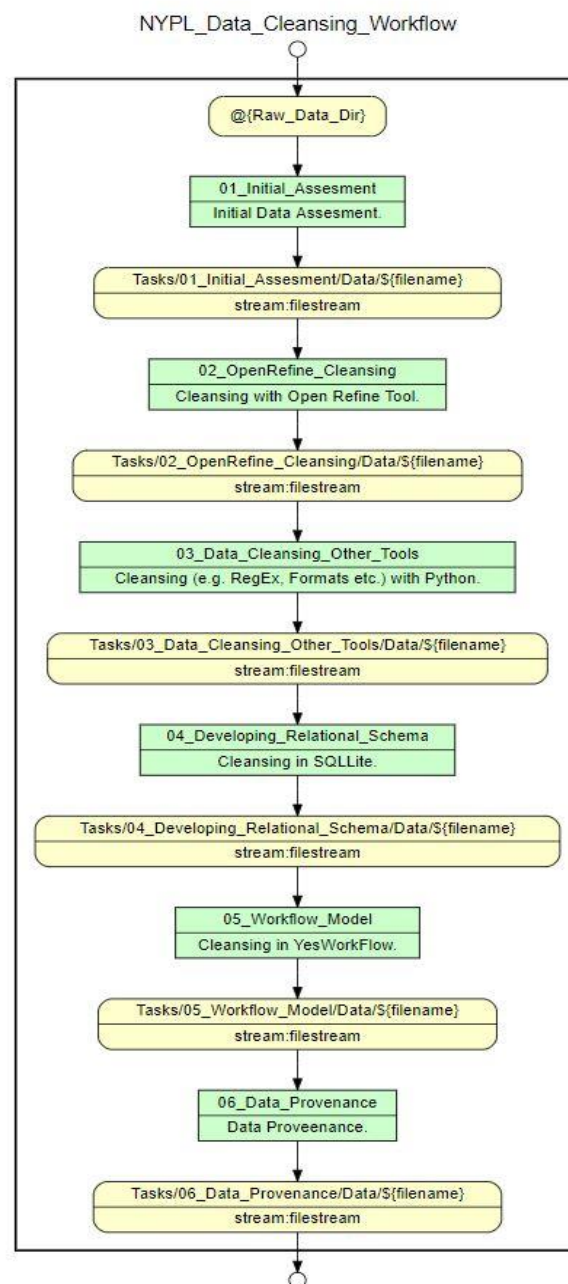
CREATE SET TABLE MenuPage
(
    Id                INTEGER NOT NULL,
    Menu_id           INTEGER NOT NULL,
    Page_number       INTEGER NOT NULL,
    Image_id          BIGINT,
    Full_height       INTEGER,
    Full_width        INTEGER,
    UUID              VARCHAR(50) CHARACTER SET LATIN NOT CASESPECIFIC,
    PRIMARY KEY (id),
    CONSTRAINT MenuPage_PK FOREIGN KEY (menu_id)
        REFERENCES Menu (id) ON UPDATE NO ACTION ON DELETE NO ACTION
);
  
```

The data was finally exported back into CSV files in the data directory of the respective task.

## 5 Creating a workflow model

The data wrangling pipeline was automated in a workflow wherein input files are transformed through individual cleaning tasks. For each of these tasks, the data lineage and processing information is retained as provenance information, so that the entire cleanup workflow can be reproduced and audited. In the following chapters, we will discuss several of these workflow components and their provenance model.

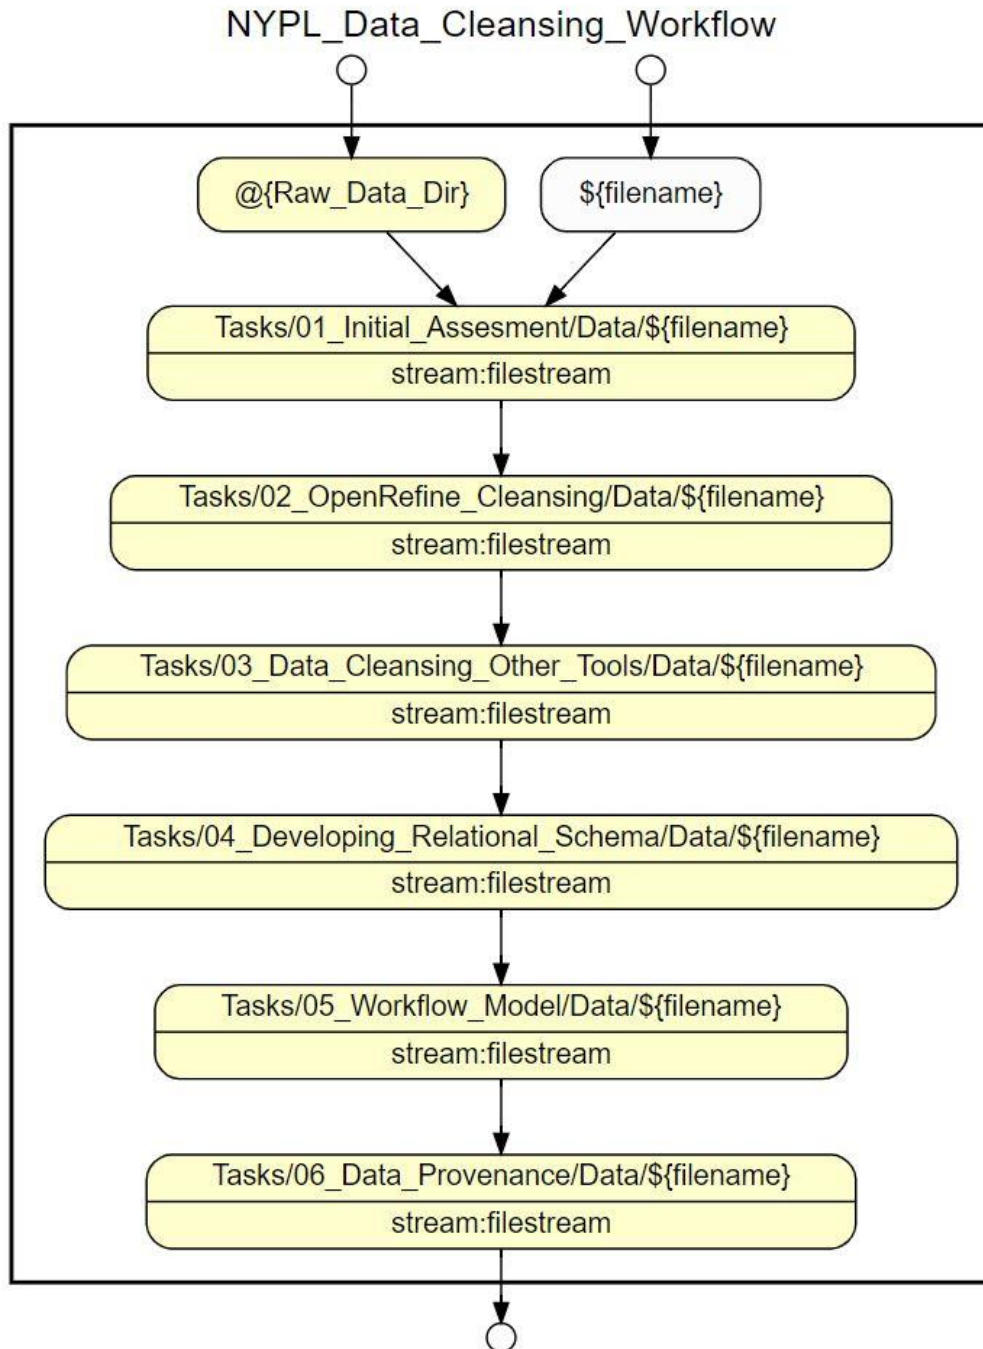
The main tool used for generating workflow provenance diagrams was YesWorkflow Online Editor. Respective workflow jobs were annotated and used to generate the respective workflow diagrams.



## 5.1 Overall Project workflow

The project data processing pipeline was set up such that each task takes the data from the preceding task's data directory, processes this data and stores it in its own data directory for the subsequent task. The first task takes the data from the project main raw data directory and the final task produces the cleansed data in its directory as the output of this entire workflow.

The main project workflow (workflow.sh) was annotated with yes workflow annotations and processes through an online editor to have the workflow graph generated for the whole project. For each individual workflow step, the flow of data from raw files to the final output file can be seen in the following graph.

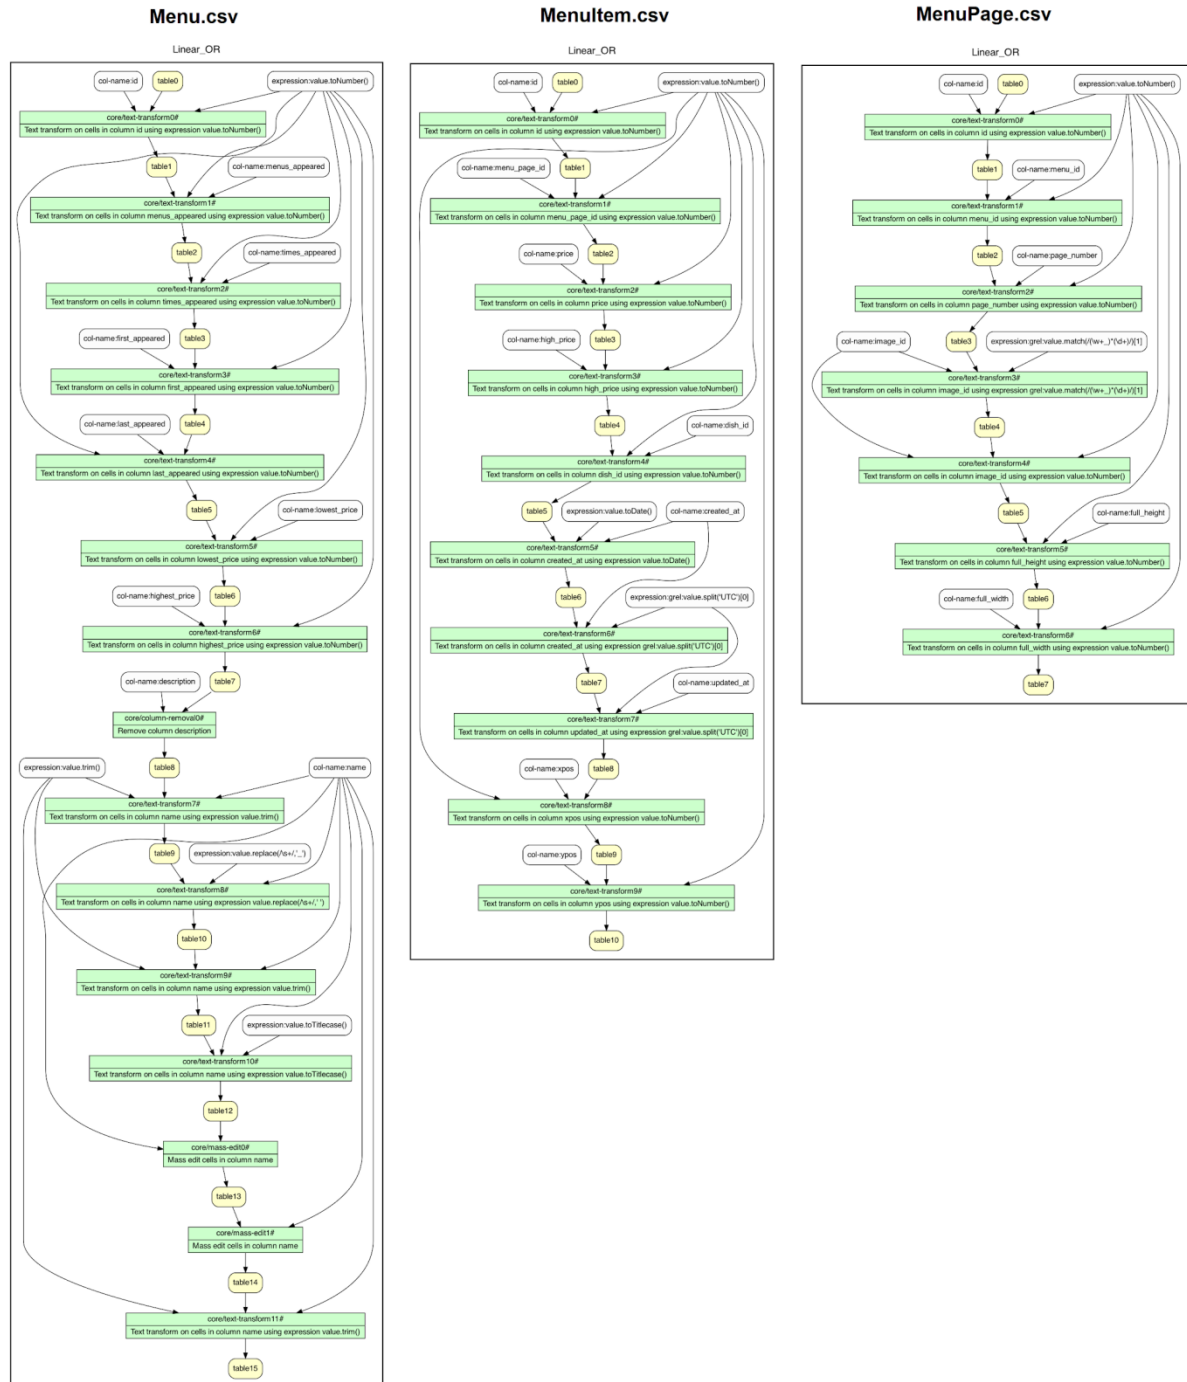


Note: To clean up the project and start fresh (delete all processed data), run **cleanup.sh**

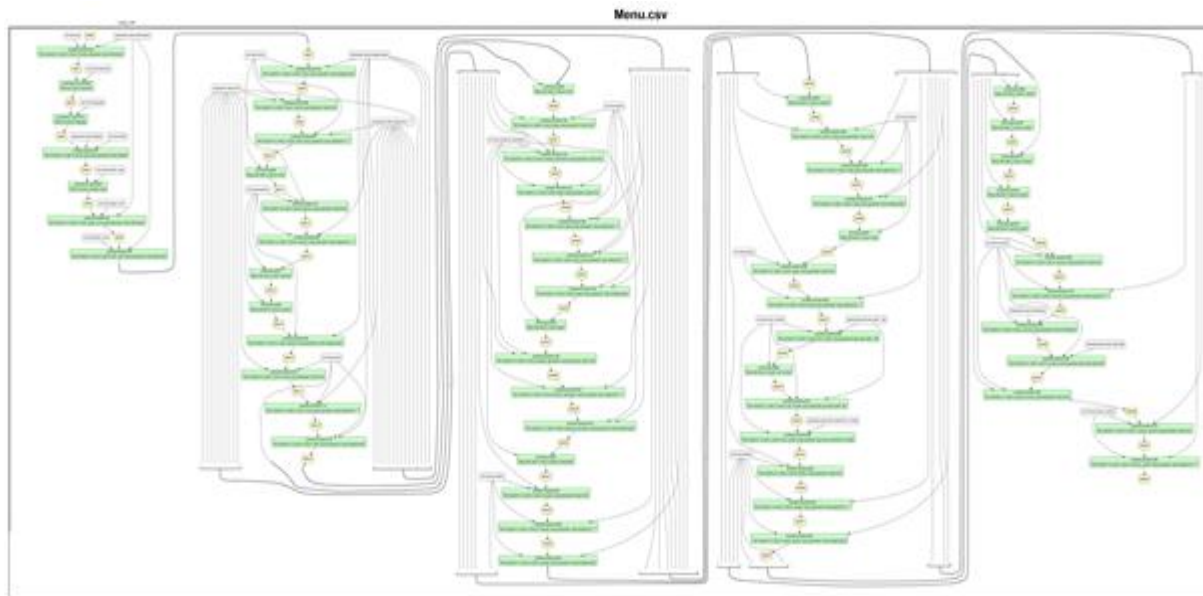
## 5.2 Open-Refine workflow

In the following are workflow graphs for each of the four input files cleansed with Open-Refine. The OR2YW tool was used to get a visual representation of the OpenRefine workflow. The JSON provenance file from OpenRefine was used to generate graphs for the workflow. The workflow was linear so changing the sequence of steps would not make much of a difference. Since the data was already in four different files in the form of a relational schema we had to clean them separately in OpenRefine. So four workflows were generated using OR2YW. The graphs for three of the input file files (**Dish.csv**, **MenuPage.csv** & **MenuItem.csv**) are as below:



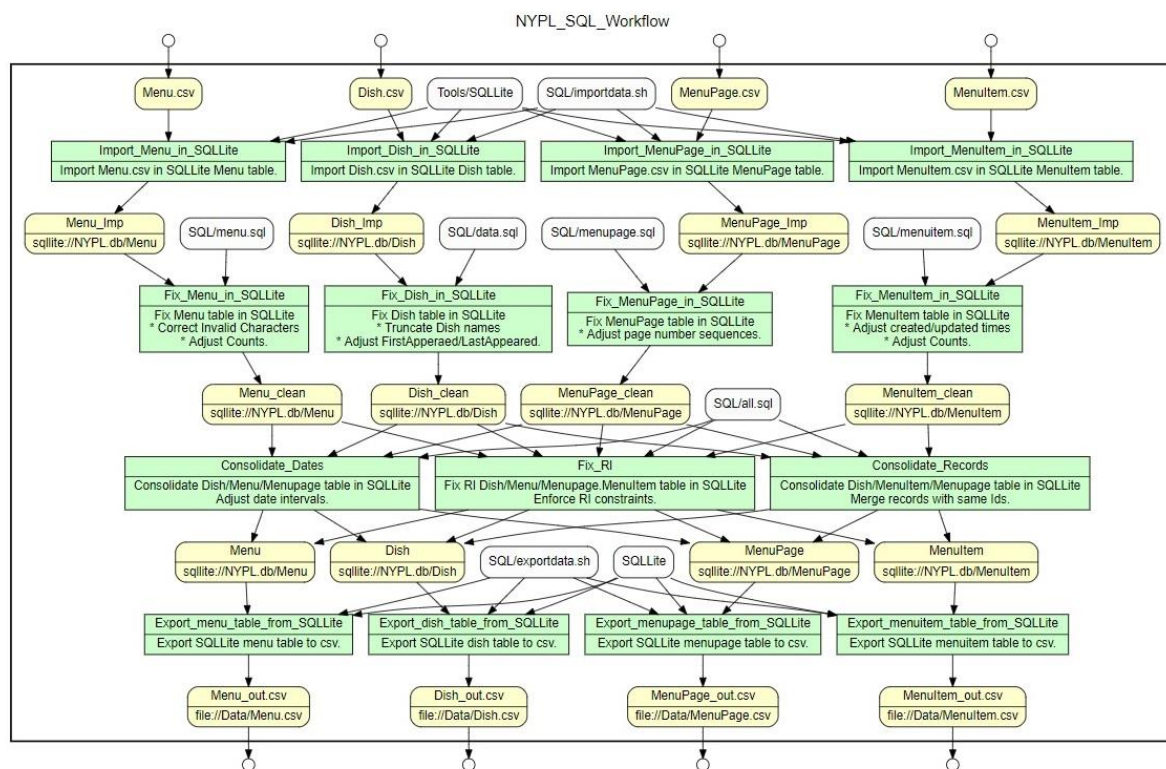


The graph of the **Menu.csv** was a bit longer and is modified to fit this report width and can be seen as below:



### 5.3 SQL-Lite workflow

Below is the work flow graph of cleansing with SQL-Lite. The input files are the pre-processed CSV file. The output files are exported CSV files after cleansing in SQL-Lite.

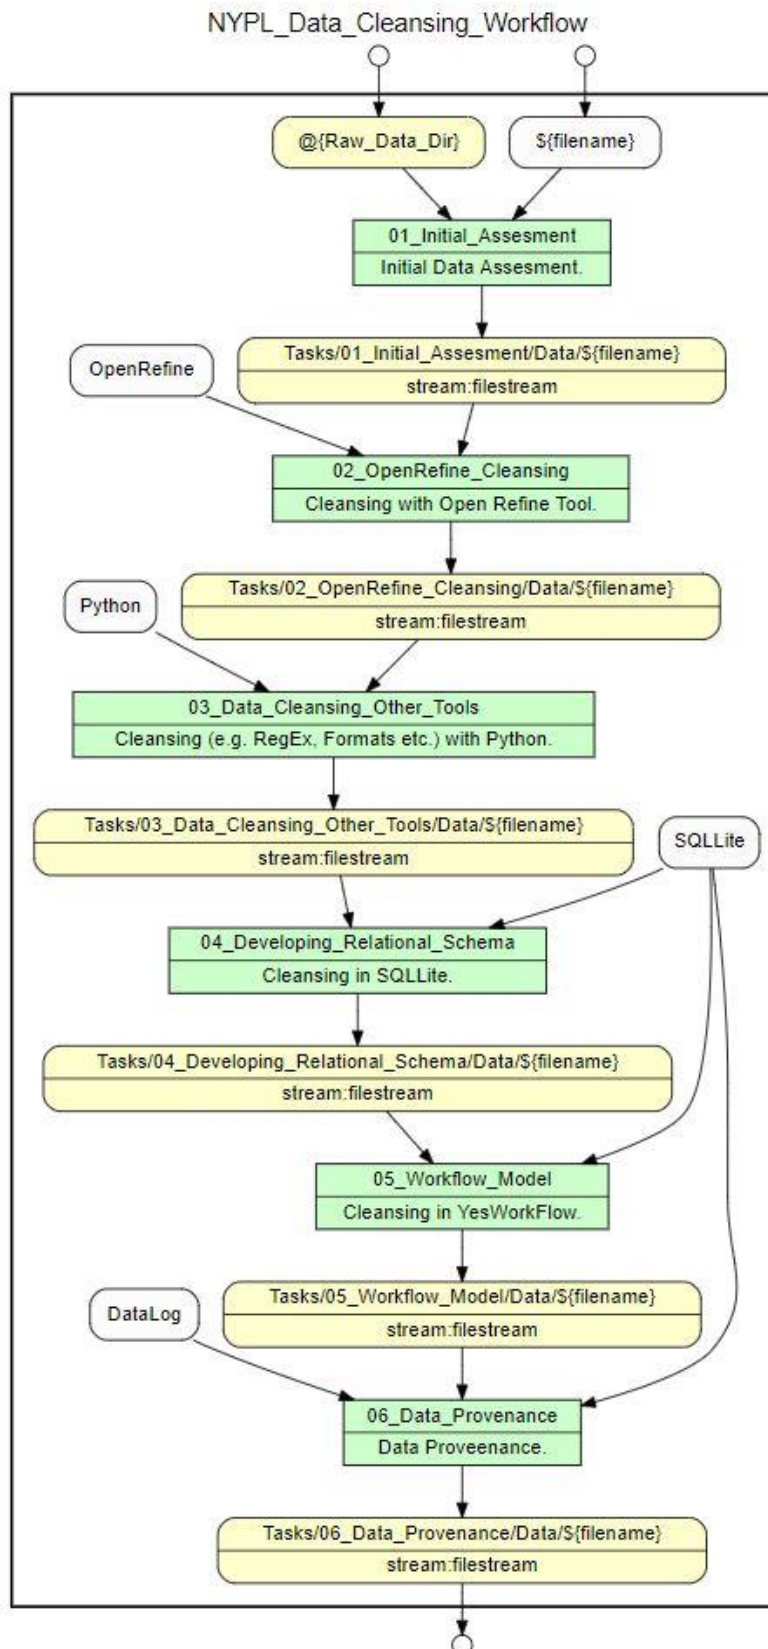


Please not that for certain updates, Teradata RDBMS was used to gain performance. The SQL scripts can be found in respective task's SQL directory.



## 5.4 Workflow inputs, Outputs and dependencies

The following workflow graph identifies key inputs, outputs, and dependencies of the overall workflow.



## 6 Developing Provenance

A provenance model was developed using clingo which is an answer set programming language used to solve combinatorial problems. Using clingo a workflow model was defined by defining input-process-output links for each task that was performed input and output being the data involved in the process. An example is shown below:

```
process(menu_file, replacenan, menu_file2).
process(menu_file2, replaceand, menu_file3).
process(menu_file3, csvtosql, nypl_db).

process(menu_page_file, fillna, menu_page_file2).
process(menu_page_file2, replaceps_rbk, menu_page_file3).
process(menu_page_file3, csvtosql, nypl_db).
```

Then a datalog rule was made to identify all the processes required to obtain a certain output file within the workflow, e.g:

```
links(X, Y, Z) :- process(X, Y, Z).
links(V, W, Z) :- process(X, Y, Z), links(V, W, X).

outputdependentprocess(Y) :- links(X, Y, menu_page_file2).
```

This query outputs the dependent process ***outputdependentprocess(fillna)*** since this output file is dependent on only one process.

Similarity another rule was made to identify all the files required to obtain a certain output file within our workflow:

```
outputdependentfiles(X) :- links(X, Y, menu_page_file2).
```

This query outputs the dependent file ***outputdependentfiles(menu\_page\_file)*** since this output file is dependent on only one file.

## 7 Conclusions and Future Work

Data wrangling is a very crucial step in data analysis and comes between data gathering and analysis. Unfortunately, this step is often ignored, and many go directly to analyze or process the data after its sourcing. The effect on analysis results can be drastic depending on whether the data was cleansed or not. Therefore, Data cleaning and wrangling is extremely effective not only in fixing data types and formats or filling in gaps but also for better analysis results.

In this project, an attempt was made to clean the NYPL "What's on the menu dataset" [4] to make more analysis-ready. For this purpose, open-source tools such as Open-Refine, SQL-Lite, and python were used. The provenance graphs were generated with Yes-Workflow Online editor, as well as with the OR2YW tool. For certain tasks for better visualization or performance, closed source tools such as DB-Visualizer and Teradata RDBMS were used, however, these are not required and used by the main project workflow.

Since several cleaning tasks were identified at the time of initial assessment which were interdependent on one other, especially in the case of integrity violations, the division of project work and time was in itself a challenge, which was managed by taking advantage of time zone difference and daily standup calls, to efficiently and effectively distribute tasks.

All workflows were created keeping the provenance of data in mind from the inception. The architecture of processing pipeline was designed such that any deliverable or data is easily traceable, reproducible and audited. A provenance model was generated in cling so that it is easier to query and identify dependencies in the workflow.

Since duplicated dishes with respect to their names were consolidated and redundancies were removed, the final dataset is much more concise and accurate and can serve multiple new use cases which was difficult if not possible in its raw form. The counts of records, respective entity prices and year of first and last appeared have been consolidated and corrected. The processed data would now be fit for analyzing how eating preferences evolved and evaluating the popularity of dishes for a given interval based on their consolidated counts or appearance. Removal of redundancies and standardization of entity names enabled easy and consistent integration of this data with another review dataset (such as Yelp reviews dataset) for rating analysis of dishes or its use in recommender systems.














In the future, the processed data could be beneficial to anyone willing to consume it for analytical work or for further wrangling tasks.

## 8 Appendix

### 1 SSH-1 GitHub Hey:

```
-----BEGIN OPENSSH PRIVATE KEY-----
b3BlbnNzaC1rZXktdjEAAAABAG5vbmUAAAABbm9uZQAAAAAAAAABAAAblwAAAAAdzc2gtcn
NhAAAAAwEAAQAAAEAviThY9kFTDex5+W4bZm1sQfDphFNeozLIFH3Q43sym94AXsnnHsu
m+2iLkClnYkKbCzR8bCTP0dQtkdbD9oJix7wkdcyM2ieoUenrSMNcVcqG/12sqrPm6SSr
jb28pBFQXj5gIUgkmdxtdbYoj8fKBx+6oMjJg0+6fepIqArYgPFFh64S8jFy6oqm6rUV
Tfvy+NQYsI4xDS4nvbbFQfzujCDjd4bh5yhQu3uhXqJ+tsWoDGphIxU1WQx2eMnnCiAZG
1zce6bL5pZfqxBpr4j4u6XJs4WZN3YIbiUleKLZPMPfCibwg9/1m7OXOVzn/yqbNdCA4pJ
TdJ5tpAr7A9IW448SUdhTdduIfVmlHe7K9Psr+KgQMRRF0NLouv4SaKo7Z69LniSILnf9T
Ni4ojQGpI32k2zju3jjBaCcet0IIYy5Jj5pz3HMM+zLfx8MWhyYSrBh9q9muls4OLnGYo6
386VODwE7BK4kDf2WxSUpZ4xcwftK8RyTZEzksbBAAAFkNaNgTHWjYExAAAAAB3NzaClyc2
EAAAGBAL4k4WPZBUw3sefluG2ZtbEHw6YRTXqMyYBR90ON7MpveAF7J5x7Lpvt0i5ApZ2J
Cmws0fGwkz9HULZHWw/aCSMe8JHXHcjNonqFH60jDXFXKhv9drKqz5ukq429vKQRUF4+
YCFIJJnXcbXW2KI/HygcfuqDIYYPun3qSIJAK2IDxRYeuEvIxcuqKpu1FU378vjUGLCO
MQ7EuJ722xUH87owg43eG4ecoULt7oV6ifrbFqAxqYSMVNVkMdnjJ5wogGYNc3Humy+aWX
6sQaa+I+LulybOFmTd2CG4lJXii2TzD3wom8IPf9Zuzlzlcl5/8qmqXQgOKSU3SebaQK+wP
SFuOPElHYU3XbiH1ZpR3uyvT7K/ieEEDURdDS6Lr+EmiqO2evS54kiC53/UzYuKI0BqSN9
pNs47t44wWgnHrdCCGMuSY+ac9xzJvSy38fDFocmEqwYfavZrtbODi5xmKot/OldA8BOWS
uJA39lsUlKwEMXMH7SvEck2RM5LGwQAAAAMBAAEAAAGAJAlGFvJtCLh+qONrhoxgYMKCHB
inwhaz5NrlsQ9V2EQuUsmRByBsA/CYkYB8ZlQ3f68310WNha/147m/0E+c7+XL4zsQUKR8
rBJVbe35r7BOEzm3Oda3hUzdVAE03oWItX26aj/2t4VRr+WmX0CT9Cnw5YGgrnPS3BwgbN
MMDYOG7UHR1xnwSa+KhGHpqBLRiUm+FePREpjDQHOYNbvEakZpteHA78MvIRVXliMa2IAV
bVThAlqnM5SP+S3YLeHdqtXk3Rep8j8ziFLRxB/sKiCMjLtwYpeMlc6L3Hae4ruoII+EVc
mXHNzxkIatE0es7G/xE+riwJ6+kYlPd+LBcDQfgy3LBkKAiz5HT1QU/vwvAVZxLFU0ShK
q078VJyTC5Tb8xG6Iq3AWgH20x6UigZPYPj4A+3Ya9lbbIOy74GzmHufev6j4sBoZr5n2
iH1t6tP/sle9PwLJ9OoulCzs0QED5N1f68glANZnIwVKgjfL3CerlcZJBRWuiYOpGBAAAA
wQcnblIV92ua8gZWITiYW6QelOYtVQxJyZkTRAFhYUDsY69ule67ZWBrHAD4WUpJYlxxC4
TG/nwRjzstjQbCi019PDcxTa6a0cCqIOVOPg4ZQnsabDKRq4loWxaiEg0ZhUsGgu1QLwnG
fF8MxfsoGF36bHy3lM2+PviI9P779EoDg6zBgP7de6CXe09xXcl1X/q3SQHf2JI4eyTfcQ
qFpzZ7ltEVmmxAK1M7jWuqML6V4iEYqJ68nKEbbU16XbE1TScAAADBA07tQ17PUGKI3jeQ
yhTfTDU2S09ooB8kLZaJludwppfHgWGNq82ATowEynxEUzviQ2PcR8NWQANAagopFPI+m0
ya7CmoLLBCizpZNCwsOVjWHXp8JteJqfVvEg9HgSJQ8QasNDO7f4JxJZXh+c8RPqQ9ILTT
WyV2khl5mQu8uWncMq7y5WomKxbeFunnvDo8lJ1jFaATBzAJ4XKqiKmt5Fdx/XravhKkP
o97Aal/ezFVxvKlnrw+kfhI57wDdLd2QAAAMEAy7s6Vz/4QEhnz2csdpSLbLx9nOB3ilpS
A0tHeEhWARoygyKz9jrthrfMG5J+htEpo9GnmWTaCyMcMzLMRGA0zP5wuqeqNNe3Q39i/q
1U7UGDfON1X3lGrE3f0TB/j4tBG9IqOxweY80ETSeR/PufeMmoaFWQ1NvSGzjj0uZ3gVlh
xfa9XvvnIuEhS8nIyUpF2wYY366XD9lcQwwyPKD6AbnNPq1vkMWvF4VYNUaCkfuC00R6nT
JzM2UPaCFu3NcpAAAFWFzYWRfQERFU0tUTlAtN1E5UE5NOAECaWQF
-----END OPENSSH PRIVATE KEY-----
```

### 2 Project Structure:

Name	Status	Date modified	Type	Size
 .git	✓	8/3/2020 8:45 AM	File folder	
 Artifacts	✓	7/10/2020 7:25 PM	File folder	
 Data	↻	8/3/2020 1:31 PM	File folder	
 Reports	↻	8/3/2020 1:58 PM	File folder	
 Tasks	✓	7/29/2020 11:59 AM	File folder	
 Tools	✓	7/16/2020 1:58 AM	File folder	
 .DS_Store		7/5/2020 11:40 AM	DS_STORE File	9 KB
 .gitignore	✓	7/16/2020 3:02 AM	Text Document	1 KB
 cleanup.sh	✓	8/3/2020 1:32 PM	Shell Script	1 KB
 Initial	✓	7/16/2020 10:49 AM	File	0 KB
 README.md	✓	7/1/2020 2:17 PM	MD File	1 KB
 workflow.sh	✓	7/29/2020 11:29 AM	Shell Script	4 KB
 workflow_annotated.sh	✓	7/29/2020 11:30 AM	Shell Script	6 KB

### 3 Main workflow script (not annotated):

```
#!/bin/bash
echo "Start of Data Cleansing"

cd Tasks

echo '#####'
echo '# Task 1 - Initial Assesment'
echo '#####'

cd "01_Initial_Assesment"

# Extract Data
echo '* Extract Data with python to: Tasks/01_Initial_Assesment/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

# Load in SQLite to Explore
echo '* Load in SQLite to Explore in: Tasks/01_Initial_Assesment/Data/NYPL.db'
./"SQL/importdata.sh" #> /dev/null 2>&1

# Execute Explore script
echo '* Explore data using: Tasks/01_Initial_Assesment/SQL/Explore.sql'
# ../../Tools/sqlite/sqlite3 Data/NYPL.db ".read SQL/Explore.sql" /dev/null 2>&1
echo '* Saving Zip Archive to: Tasks/01_Initial_Assesment/Data'

cd ..

echo '#####'
echo '# Task 2 - Open Refine'
echo '#####'

cd "02_OpenRefine_Cleansing"

# Extract Open Refine Data
echo '* Extract Data with python to: Tasks/02_OpenRefine_Cleansing/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

# Generate OpenRefine Audit
echo '* Extract OpenRefine Audit logs to : Tasks/02_OpenRefine_Cleansing/Open Refine'
echo '    OpenRefine < Data/Dish.csv > "Open Refine/Dish.json"'
echo '    OpenRefine < Data/MenuItem.csv > "Open Refine/MenuItem.json"'
echo '    OpenRefine < Data/MenuPage.csv > "Open Refine/MenuPage.json"'
echo '    OpenRefine < Data/Menu.csv > "Open Refine/Menu.json"'
echo '* Saving Zip Archive to: Tasks/02_OpenRefine_Cleansing/Data'

cd ..

echo '#####'
echo '# Task 3 - Other Tools'
echo '#####'

cd "03_Data_Cleansing_Other_Tools"

echo '* Extract Open Refine Cleansed Data with python to: Tasks/03_Data_Cleansing_Other_Tools/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

echo '    Processing Menu.csv'
python "Python/Menu.py" #> /dev/null 2>&1

echo '    Processing Dish.csv'
python "Python/Dish.py" #> /dev/null 2>&1

echo '    Processing MenuItem.csv'
python "Python/MenuItem.py" #> /dev/null 2>&1

echo '    Processing MenuPage.csv'
python "Python/MenuPage.py" #> /dev/null 2>&1

echo '* Saving Zip Archive to: Tasks/03_Data_Cleansing_Other_Tools/Data'
python "Python/ZipTaskData.py" > /dev/null 2>&1

cd ..

echo '#####'
echo '# Task 4 - Relational Schema'
echo '#####'

cd "04_Developing_Relational_Schema"

echo '* Extract Open Refine Cleansed Data with python to: Tasks/04_Developing_Relational_Schema/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

# Load in SQLite to Explore
echo '* Load in SQLite to Explore in: Tasks/04_Developing_Relational_Schema'
./"SQL/importdata.sh"
echo '* Saving Zip Archive to: Tasks/04_Developing_Relational_Schema/Data'
#python "Python/ZipTaskData.py" > /dev/null 2>&1

cd ..

echo '#####'
echo '# Task 5 - Workflow Model'
echo '#####'

cd "05_Workflow_Model"

echo '* Extract Open Refine Cleansed Data with python to: Tasks/05_Workflow_Model/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

echo '* Saving Zip Archive to: Tasks/05_Workflow_Model/Data'
python "Python/ZipTaskData.py" > /dev/null 2>&1

cd ..

echo '#####'
echo '# Task 6 - Data Provenance'
echo '#####'

cd "06_Data_Provenance"

echo '* Extract Open Refine Cleansed Data with python to: Tasks/06_Data_Provenance/Data'
python "Python/ExtractData.py" > /dev/null 2>&1

echo '* Saving Zip Archive to: Tasks/06_Data_Provenance/Data'
python "Python/ZipTaskData.py" > /dev/null 2>&1

cd ..

echo "End of Data Cleansing"

cd ..
```

#### 4 Workflow log:

```
(base) C:\>sh workflow.sh
Start of Data Cleansing
#####
# Task 1 - Initial Assesment
#####
* Extract Data with python to: Tasks/01_Initial_Assesment/Data
* Load in SQLite to Explore in: Tasks/01_Initial_Assesment/Data/NYPL.db
" Dish file added, Total Dishes = 426985"
" Menu file added, Total Menus = 17545"
" MenuPage file added, Total MenuPages = 66937"
" MenuItem file added, Total MenuItem = 1334419"
* Explore data using: Tasks/01_Initial_Assesment/SQL/Explore.sql
* Saving Zip Archive to: Tasks/01_Initial_Assesment/Data
#####
# Task 2 - Open Refine
#####
* Extract Data with python to: Tasks/02_OpenRefine_Cleansing/Data
* Extract OpenRefine Audit logs to : Tasks/02_OpenRefine_Cleansing/Open Refine
  OpenRefine < Data/Dish.csv > "Open Refine/Dish.json"
  OpenRefine < Data/MenuItem.csv > "Open Refine/MenuItem.json"
  OpenRefine < Data/MenuPage.csv > "Open Refine/MenuPage.json"
  OpenRefine < Data/Menu.csv > "Open Refine/Menu.json"
* Saving Zip Archive to: Tasks/02_OpenRefine_Cleansing/Data
#####
# Task 3 - Other Tools
#####
* Extract Open Refine Cleansed Data with python to: Tasks/03_Data_Cleansing_Other_Tools/Data
  Processing Menu.csv
  Processing Dish.csv
  Processing MenuItem.csv
  Processing MenuPage.csv
* Saving Zip Archive to: Tasks/03_Data_Cleansing_Other_Tools/Data
#####
# Task 4 - Relational Schema
#####
* Extract Open Refine Cleansed Data with python to: Tasks/04_Developing_Relational_Schema/Data
* Load in SQLite to Explore in: Tasks/04_Developing_Relational_Schema
  Truncating Dish Names Longer than 500 Characters"
" Correct Dishes where last_appeared < first_appeared: 6 rows affected"
" Correct Dishes where last_appeared > Current Date : 179 rows affected"
" Correct Dishes where first_appeared > Current Date : 11 rows affected"
" * Dish table corrected"
" Correct Invalid characters in Menu : 179 rows affected"
" * Menu Table Corrected"
" MenuPage.pagenumbers resequenced : 175 rows affected"
" * MenuPage Table Corrected"
" Correcting price > high_price : 2378 rows affected"
" Correcting created_at > updated_at : 2874 rows affected"
" * MenuItem Table Corrected"
" Consolidating dishes : 36701 rows affected"
" Consolidating MenuItem : 23371 rows affected"
" Fixing RI in MenuPage"
" Updating Dish times appeared : 21693 rows affected"
" Updating Dish menus appeared : 28843 rows affected"
" Updating Dish first appeared : 25272 rows affected"
" Updating Dish last appeared : 17857 rows affected"
" Updating Dish prices : 20897 rows affected"
" Updating Dish counts : 82967 rows affected"
" Updating Dish Page counts : 214 rows affected"
" * All Data Corrected with SQL-Lite"
" * SQL-Lite DB Exported"
* Saving Zip Archive to: Tasks/04_Developing_Relational_Schema/Data
#####
# Task 5 - Workflow Model
#####
* Extract Open Refine Cleansed Data with python to: Tasks/05_Workflow_Model/Data
* Saving Zip Archive to: Tasks/05_Workflow_Model/Data
#####
# Task 6 - Data Provenance
#####
* Extract Open Refine Cleansed Data with python to: Tasks/06_Data_Provenance/Data
* Saving Zip Archive to: Tasks/06_Data_Provenance/Data
End of Data Cleansing
```



## Bibliography

- [1] David Huynh, Stefano Mazzocchi, Metaweb Technologies, Inc, "OpenRefine," October 2012. [Online]. Available: <https://openrefine.org/>.
- [2] SQLite Consortium, "SQLite," [Online]. Available: <https://www.sqlite.org/index.html>.
- [3] DBVis Software, "DbVisualizer: A universal database tool," DbVis Software AB, [Online]. Available: <https://www.dbvis.com/>.
- [4] NYPL Labs, "What's on the menu Dataset," 16 June 2020. [Online]. Available: <http://menus.nypl.org/data>.
- [5] NYPL Labs, "Whats-On-The-Menu," June 2020. [Online]. Available: <http://nypl.github.io/menus-api/>.

## Contributors:



**Asad Bin Imtiaz** [Net ID: **aimtiaz2**] worked on:

- 1> Initial assessment (Task 1) in identifying the potential issues, writing queries to identify issues and preparing directory structure
- 2> Data Cleansing with Open Refine (Task 2) in clustering records and case standardization
- 3> Data Cleansing with Other tools (Task3) in writing python scripts
- 4> Developing relational schema (Task 4) in writing SQL queries in Teradata and developing SQL Schema
- 5> Creating workflow model (Task 5) in creating overall workflow and directory structure, created workflow for SQLs and generating provenance graphs
- 6> Writing of this project report.



**Mohammad Rafay** [Net ID: **mrafay2**] worked on:

- 1> Initial assessment (Task 1) in identifying the defining use cases, data set structure and data quality assessment.
- 2> Data Cleansing with Open Refine (Task 2) in fixing data type and standardization issues.
- 3> Developing relational schema (Task 4) in writing SQL queries in SQL Lite and fixing RI issues.
- 4> Creating workflow model (Task 5) in generating workflow for OpenRefine logs and their respective provenance graphs.
- 5> Developing provenance (Task 6) in developing data log statements using clingo.
- 6> Writing of this project report.